

Universitat Politècnica de Catalunya
Facultat d'Informàtica de Barcelona

COMPUTING & INTELLIGENT SYSTEMS

Lab 3 - Machine Learning

Robert CARAUSU
Marc VILA GÓMEZ

January 17, 2018

Contents

| | | |
|----------|--|-----------|
| 1 | Introduction | 1 |
| 1.1 | Problem statement | 1 |
| 1.2 | Summary | 2 |
| 2 | Data analysis & preprocessing | 3 |
| 2.1 | Continuous variables | 3 |
| 2.2 | Categorical variables | 5 |
| 3 | Modelling | 7 |
| 3.1 | Logistic Regression | 7 |
| 3.2 | LDA, QDA & Naive Bayes | 10 |
| 3.3 | Random Forest | 10 |
| 4 | Conclusions | 12 |
| 5 | Annex | 13 |
| 5.1 | Code | 13 |

1 Introduction

1.1 Problem statement

A Portuguese banking institution has made a **direct marketing** campaign (by means of phone calls) in order to make clients subscribe to certain financial product (a term deposit). The data has the following characteristics:

- Often, more than one contact to the same client was required, in order to assess if the product would be subscribed.
- Number of examples: 45,211; 16 predictors, of very different nature and type, including factors, '999' and 'unknown'
- The **target** variable is whether a term deposit was subscribed ('yes') or not ('no')

The **input** (predictive variables) are:

- bank client data:
 1. age (numeric)
 2. job : type of job ("admin.", "unknown", "unemployed", "management", "house-maid", "entrepreneur", "student", "blue-collar", "self-employed", "retired", "technician", "services")
 3. marital : marital status ("married", "divorced", "single")
 4. education ("unknown", "secondary", "primary", "tertiary")
 5. default: has credit in default? ("yes", "no")
 6. balance: average yearly balance, in euros (numeric)
 7. housing: has housing loan? ("yes", "no")
 8. loan: has personal loan? ("yes", "no")
- related with the last contact of the current campaign:
 1. contact: contact communication type ("unknown", "telephone", "cellular")
 2. day: last contact day of the month (numeric)
 3. month: last contact month of year ("jan", "feb", "mar", ..., "nov", "dec")
 4. duration: last contact duration, in seconds (numeric)
- other variables:
 1. campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)

2. pdays: number of days passed after the client was last contacted from a previous campaign (numeric, -1 means client was not previously contacted)
3. previous: number of contacts performed before this campaign and for this client (numeric)
4. poutcome: outcome of the previous marketing campaign ("unknown", "other", "failure", "success")

The bank institution wants to predict the probability that a client will subscribe or no the deposit based on the previous data, so they segment the market and center their resources on the clients with a higher probability of subscribing it.

1.2 Summary

In this document we will perform the analysis and selection of the best classification method in order to solve the problem stated in the previous section, which is to correctly predict the class of the target variable *deposit subscribed* (**yes**), or not (**no**).

We will first perform an analysis of the data that we have available from the bank and do any processing that may be necessary to it in order to increase the efficiency of the methods that we'll test.

Then, we will test five different classifiers: **Logistic Regression**, **LDA**, **QDA**, **Naive Bayes** and **Random Forest**, and select the best one that fits our model based on its **error** and **precision**. The last one, **Random Forest**, is especially important since as we will find out the data is highly unbalanced, with just 11.70% of the clients subscribing to the deposit after the marketing campaign.

2 Data analysis & preprocessing

In this section we will discuss and analyze the data that was given to the students and also perform some preprocessing that may be necessary in order to perform the modelling.

2.1 Continuous variables

In Figure 1 we can see the continuous variables from the data before processing. The age variable seems to be ok, but the other variables are highly skewed and also need to be scaled. In order to fix these issues we'll be applying the *log* and *scale* functions wherever we can:

- balance: it has negative values so we can't apply log to fix the skewness without removing data, so we will just scale it.
- duration: we can fix both skewness and scale.
- campaign : if we apply scale and log it has some undesired effects (as seen in Figure 2) so we leave it as it is.
- previous: if we apply scale and log it has some undesired effects (as seen in Figure 2) so we leave it as it is.
- pdays: most of its values are -1 (which means that the client was not contacted in a previous campaign) so we decide to create a categorical variable with *not_contacted* and *contacted* as categories.

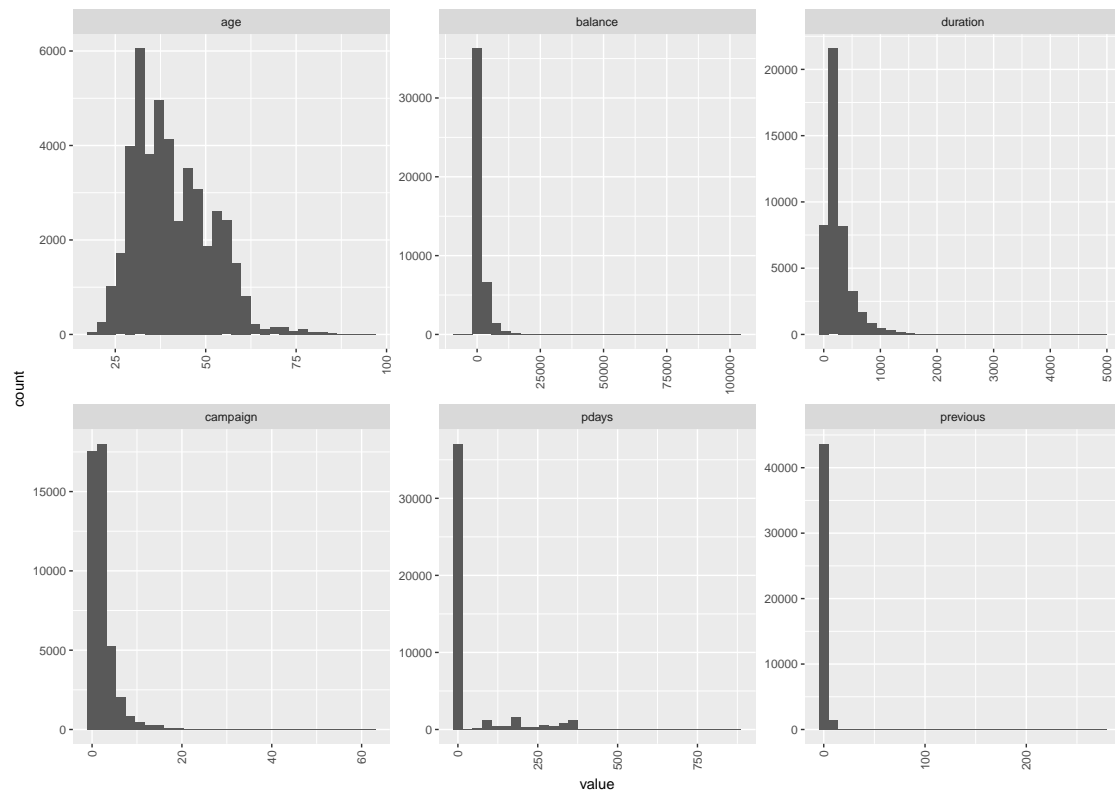


Figure 1: Continuous variables before preprocessing

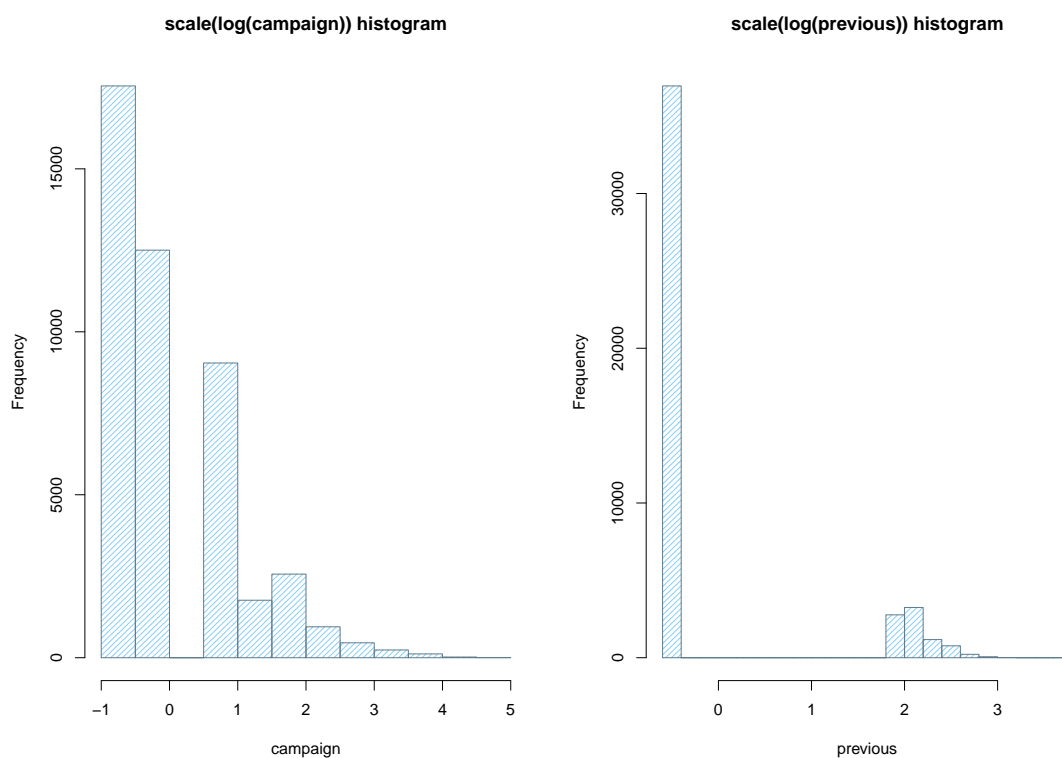


Figure 2: Applying log and scale to campaign and previous variables

After applying all the changes mentioned in the previous paragraphs, the continuous variables remain as seen in Figure 3. The *pdays* variable is shown in Figure 4 with the rest of the categorical variables.

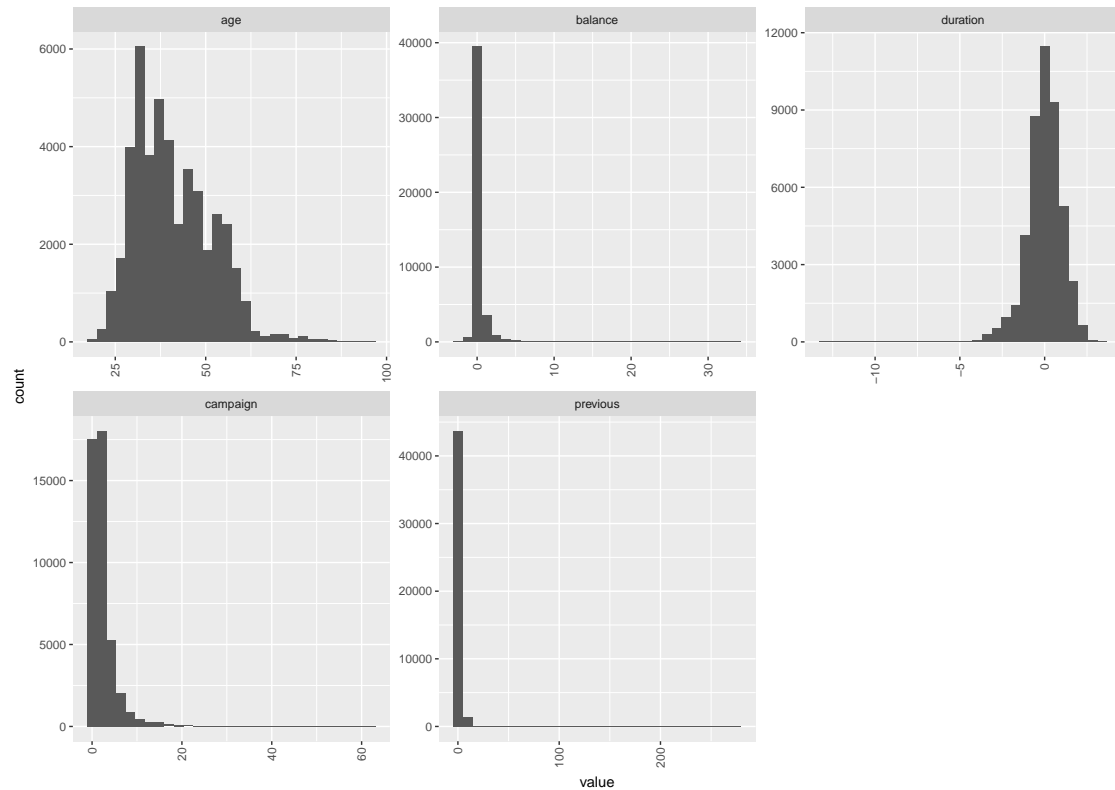


Figure 3: Continuous variables after preprocessing

2.2 Categorical variables

In Figure 4 we can see the categorical variables. As we can see, most of the variables seem ok except for *job*, *education*, *contact*, and *poutcome*, which the missing observations are labeled as *unknown*. Since we don't want to remove any of the data we will leave the categorical variables as they are. We could replace them for the **NA** value in R, but since only Random Forest can deal with this kind of variables we choose to not modify anything.

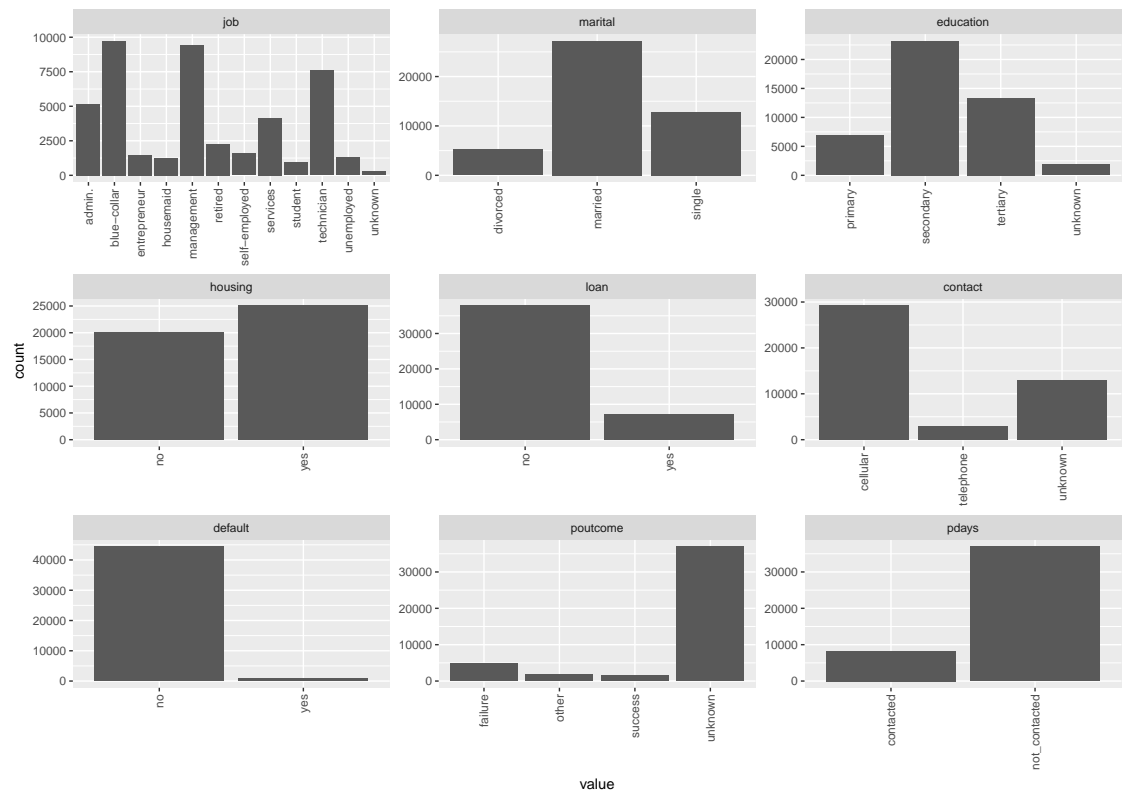


Figure 4: Categorical variables

3 Modelling

In this section we will be discussing the different classifiers that we will use to fit our data and make predictions. We will use five different classifiers as stated before: Logistic Regression, LDA, QDA, Naive Bayes and Random Forest. We will compare them by two different parameters, **error** and **precision**, this last one being of special importance since the bank is more interested in the people that may want to subscribe to their deposit, so the resources of the marketing campaign can be dedicated to contacting these clients.

Because of the nature of the data, we expect Logistic Regression and Random Forest to be the ones that perform better. In order to fit our models, we will be dividing our data in two different sets: **learning** and **testing**, the first one being used to fit the models and the second one to test it and make predictions, in a 2/3 and 1/3 proportion, respectively.

The details of the implementation can be seen in the code provided in the annex or in the `.R` file included in the folder with this document.

3.1 Logistic Regression

After fitting the first model with the original data, we inspect the *p-values* and observe that there are a few variables that have a really low impact in the target variable (the higher the *p-value* the lower the impact), these are: *age*, *job*, *marital*, *default*, *balance*, *pdays* and *previous*.

So we decide to do four different fits:

- Original data, with all the variables.
- Original data, without the variables mentioned.
- Pre-processed data, with all the variables.
- Pre-processed data, without the variables mentioned.

After fitting the first model, we find that the error is **9.93%** and the precision is **34.15%**. This is a pretty low precision, it makes us miss most of the clients that would said yes to subscribing the deposit. In order to improve this value and balance the error we lower the threshold for a positive subscription from 0.5 to 0.3, now the error is **9.97%** (a very small increase), but the precision is **53.78%**, almost doubled, so we decide to make all the fits with this new value for the threshold.

The results can be observed in Table 1, we see that the fact that we fit without the discarded variables doesn't improve the error and precision of the model (in fact it seems to lower it), so we will be using all of the variables from now on. But there is an important improvement in using the pre-processed data, we can see that the precision has improved by more than 4 points, which is quite good.

| Variables | Data | Error | Precision |
|------------------|---------------|--------------|------------------|
| All | Original | 9.97% | 53.78% |
| Less | Original | 10.07% | 53.32% |
| All | Pre-processed | 10.60% | 57.95% |
| Less | Pre-processed | 10.53% | 57.44% |

Table 1: Error and precision for logistic regression

In order to get a better grasp of the performance of the classifier we performed *k-fold Cross Validation* dividing the data in 10 equal divisions, using 1 for validation and 9 for training, and performing it 100 times since we observed that the computation time wasn't too high. The results can be observed in Table 2 which confirms our previous suppositions. In Figure 5 and Figure 6 there are the histograms and box plots for the accuracy, error and precision for logistic regression with all the variables after the cross validation

| Variables | Data | Error | Precision |
|------------------|---------------|--------------|------------------|
| All | Pre-processed | 10.31% | 58.80% |
| Less | Pre-processed | 10.26% | 58.37% |

Table 2: Mean values for error and precision for logistic regression with cross validation

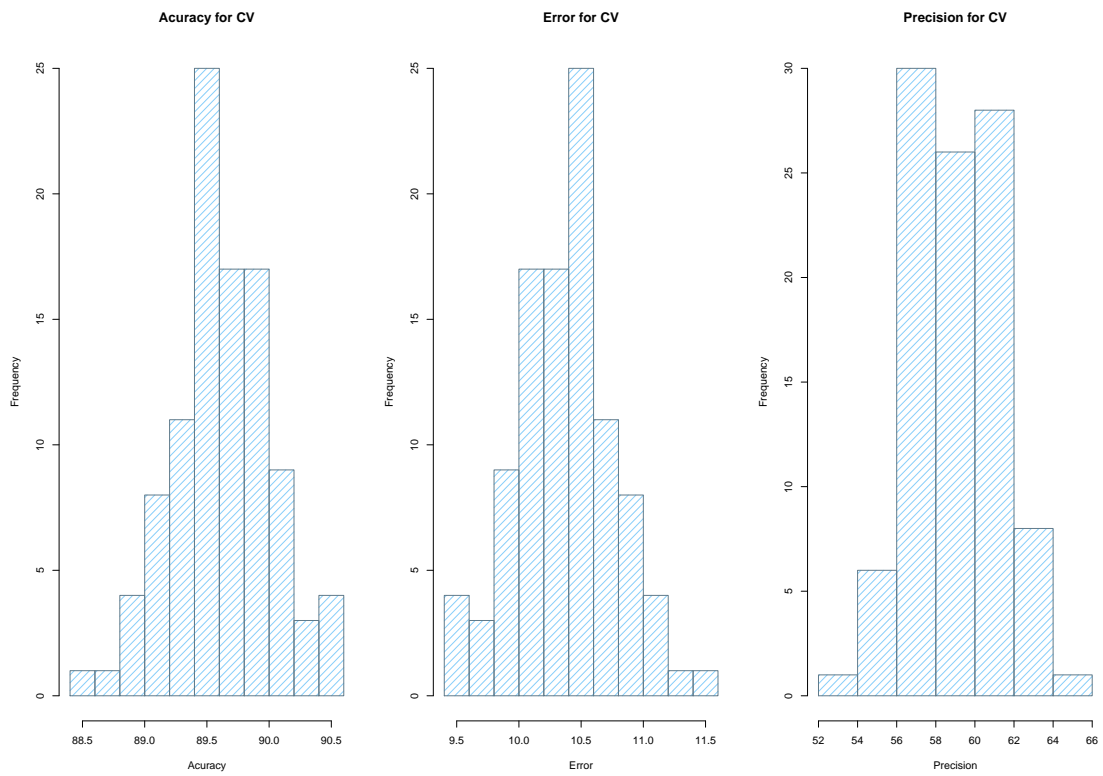


Figure 5: Histograms of accuracy, error and precision for logistic regression

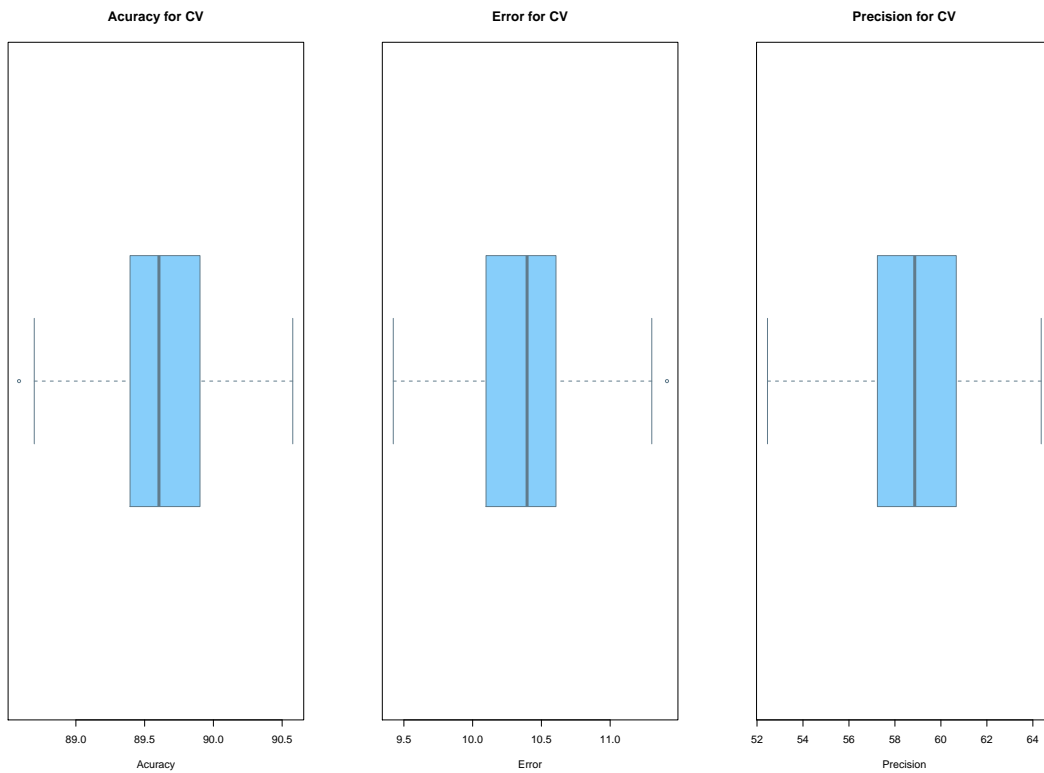


Figure 6: Box plots of accuracy, error and precision for logistic regression

3.2 LDA, QDA & Naive Bayes

We also performed the fitting for three other classifiers, LDA, QDA and Naive Bayes. The results can be observed in Table 3. We can see that as predicted, LDA and QDA performed rather poorly, both with a precision way lower than 50%, the limit with which consider a prediction better than randomly assigning the observations to subscribed or not subscribed. The only one that got close to this value is Naive Bayes but it still has a high error, so we consider these three classifiers not good for our model.

| Model | Error | Precision |
|-------------|--------|-----------|
| LDA | 10.73% | 28.51% |
| QDA | 13.43% | 37.99% |
| Naive Bayes | 12.57% | 43.89% |

Table 3: Results for LDA, QDA and Naive Bayes

3.3 Random Forest

The last classifier that we tested is Random Forest. The results of it can be observed in Table 4. Because of the nature of this classifier, we used the original untreated data with a number of trees of 300, and we got the results in the first row of the table mentioned before. So far this is the classifier with the best results, which is the one that we will decide for, since it has the lowest error and highest precision. We decided to do a 10-fold Cross Validation to make sure the values we were getting were consistent, this wasn't really necessary because of how this classifier works, but the results we got were pretty much the same as the first iteration.

In Figure 7 and Figure 8 we can observe the histograms and box plots for the accuracy, error and precision for this method. **Because of the better results of random forest, we would decide for this method to model the data for the marketing campaign of the bank.**

| Model | Error | Precision |
|----------------------------|-------|-----------|
| Random Forest | 9.30% | 64.66% |
| Random Forest (10-fold CV) | 9.17% | 64.66% |

Table 4: Results for Random Forest classifier

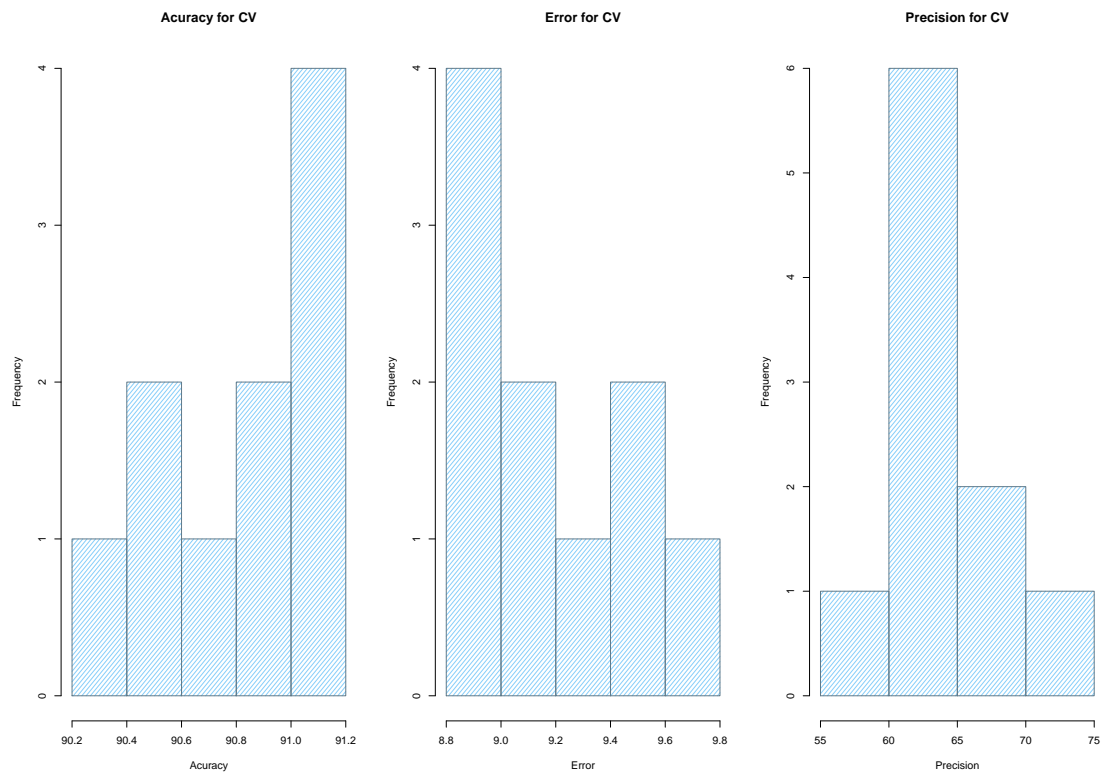


Figure 7: Histograms of accuracy, error and precision for random forest

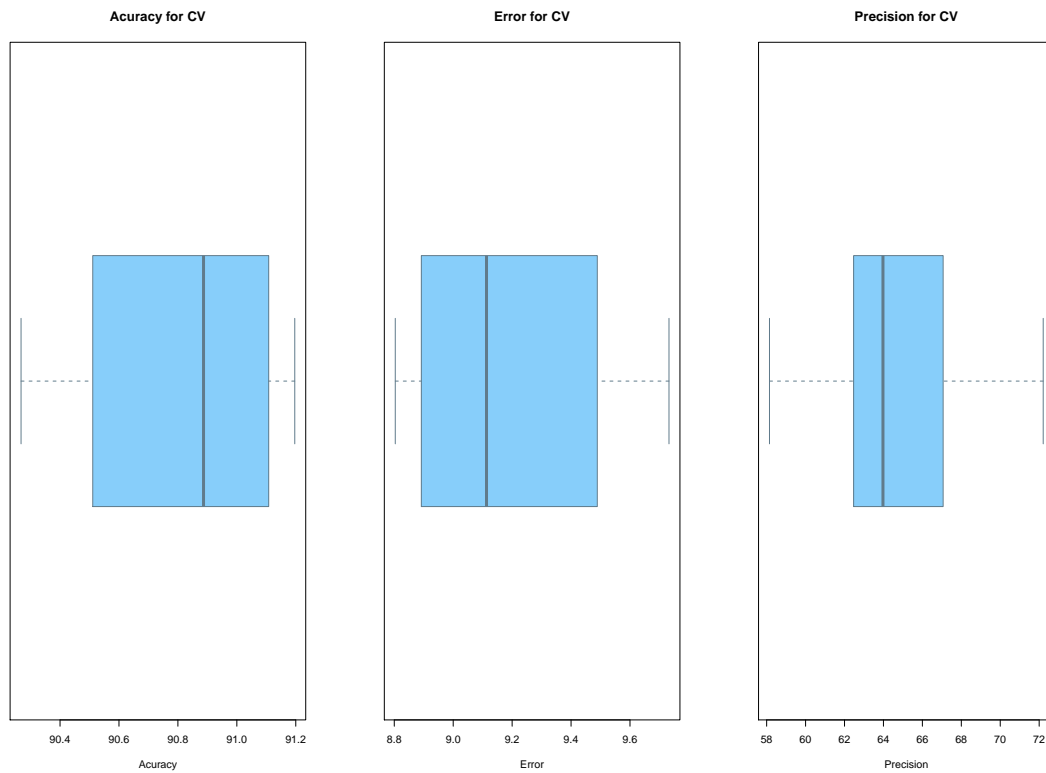


Figure 8: Box plots of accuracy, error and precision for random forest

4 Conclusions

In this document we made the analysis and comparison of five different classifiers for the problem stated in the introduction. We got the opportunity to observe how important is the preparation of the data in order to get the best results for our model, and also how choosing the appropriate classifier can mean success or failure in addressing the problem correctly.

In this particular case, we observed that the best method was Random Forest. This is because we were working with a particular set of data, which was very unbalanced and has a lot of different variables that make it more difficult to find a correlation between them and the target variable.

In general we are satisfied with the results obtained, since a precision of more than 60% means that the bank is not making random guessing in what type of client to invest time in selling its products. This values of course can probably be improved with more dedication and knowledge from our part.

5 Annex

5.1 Code

```
1 #####
2 ##### Project 1: the Bank Marketing Data Set
3 ##### Robert Carausu & Marc Vila, CSI - MEI 2017-2018
4 #####
5 set.seed (6046)
6 library(reshape2)
7 library(ggplot2)
8
9 ## Direct marketing campaigns (phone calls) of a Portuguese banking
   ↪ institution.
10 ## The classification goal is to predict if the client will subscribe a
   ↪ term deposit
11
12 ## Getting the dataset
13 deposit <- read.table(file="./data/bank.csv", header=TRUE,
   ↪ stringsAsFactors=TRUE, sep=";")
14 # We rename the target variable
15 colnames(deposit)[ncol(deposit)] <- "subscribed"
16 original_data = deposit # We make a copy to compare it later with our
   ↪ pre-processed data
17
18 # 45211 observations and 17 different variables
19 # (9 categorical: job, marital, education, default, housing, loan,
   ↪ contact, month, poutcome and y)
20 dim(deposit)
21 summary(deposit)
22 # 11.70% of subscribed, so our model sholdn't have a higher error than
   ↪ this
23 # Data is very unbalanced so some models will adjust worse than others
24 sum(deposit$subscribed=="yes")/sum(length(deposit$subscribed))*100
25
26 ## Let's have a visual inspection of the continuous variables before
   ↪ pre-processing
27 # Age seems ok
```

```

28 # The other variables are highly skewed so we will try to scale and apply
    ↪ log where we can
29 # We can do it for duration, not for balance since it has negative values
    ↪ and we don't want to lose data
30 # pdays can be converted to categorical: "not_contacted" (in previous
    ↪ campaign) and "contacted"
31 d.cont <- melt(deposit[, c("age", "balance", "duration", "campaign",
    ↪ "pdays", "previous")])
32 ggplot(d.cont, aes(x = value)) + facet_wrap(~variable, scales = "free") +
    ↪ geom_histogram() + theme(axis.text.x=element_text(angle=90, hjust=1,
    ↪ vjust=0.5))
33
34 ## Let's have a visual inspection of the factor variables before
    ↪ pre-processing
35 # They seem ok so we won't be touching these variables
36 d.categ <- melt(deposit, measure.vars=c("job", "marital", "education",
    ↪ "housing", "loan", "contact", "default", "poutcome"))
37 ggplot(d.categ, aes(x = value)) + facet_wrap(~variable, scales = "free") +
    ↪ geom_bar() + theme(axis.text.x=element_text(angle=90, hjust=1,
    ↪ vjust=0.5))
38
39 # This dataset needs a lot of pre-processing ... also it displays a good
    ↪ mixture of categorical and numeric variables
40 # In conclusion LDA/QDA may not be the best choice, a good choice may be
    ↪ Logistic Regression. We will test also Naive Bayes and Random Forest
    ↪ and
41 # choose the best model that fits our problem
42
43 ##### PRE-PROCESSING #####
44 ##### Fixing skewness and scaling continuous variables
45 # The balance has negative values, so we can only scale it
46 hist(deposit$balance, col='lightskyblue', border='lightskyblue4',
    ↪ xlab='balance', main='balance histogram', density=50)
47 # There are 3766 for negative balance
48 # The only way to fix it is to delete these observations so we choose to
    ↪ leave it as it is since we don't want to lose data
49 sum(deposit$balance<0)
50 deposit$balance = scale(deposit$balance)

```



```

51 # Scaled balance
52 hist(scale(deposit$balance), col='lightskyblue', border='lightskyblue4',
    ↪ xlab='balance', main='balance histogram', density=50)
53
54 # duration, campaign and previous are all skewed, so we apply log and
    ↪ scale
55 hist(deposit$duration, col='lightskyblue', border='lightskyblue4',
    ↪ xlab='duration', main='duration histogram', density=50)
56 deposit$duration = log(deposit$duration+0.001) # +0.001 to avoid -Inf
57 hist(deposit$duration, col='lightskyblue', border='lightskyblue4',
    ↪ xlab='duration', main='duration histogram', density=50)
58 deposit$duration = scale(deposit$duration)
59 hist(deposit$duration, col='lightskyblue', border='lightskyblue4',
    ↪ xlab='duration', main='duration histogram', density=50)
60
61 # Applying log and scale to campaign and previous has some undesired
    ↪ effects, so we will leave them as they are
62 hist(scale(log(deposit$campaign + 0.001)), col='lightskyblue',
    ↪ border='lightskyblue4', xlab='campaign', main='scale(log(campaign))
    ↪ histogram', density=50)
63 hist(scale(log(deposit$previous + 0.001)), col='lightskyblue',
    ↪ border='lightskyblue4', xlab='previous', main='scale(log(previous))
    ↪ histogram', density=50)
64
65 # pdays has most of values -1 (not contacted previously).
66 # We make a categorical value with "contacted" for pdays!=-1 and "not
    ↪ contacted" previously for pdays=-1
67 hist(deposit$pdays, col='lightskyblue', border='lightskyblue4',
    ↪ xlab='pdays', main='pdays histogram', density=50)
68 deposit$pdays = cut(deposit$pdays, breaks=c(-Inf, 0.0, Inf),
    ↪ labels=c("not_contacted", "contacted"))
69 table(deposit$pdays)
70 plot(deposit$pdays)
71
72 ##### Fixing "unknown" values
73
74 # There are 288 subscriptions for unknown job, we leave it as it is since
    ↪ we don't want to delete this data

```

```

75 summary(deposit[deposit$job=="unknown",])
76
77 # We could change the unknown values to NA (as well as the 0 previous
  ↪ contacts variables), this is useful if we use a Random Forest
  ↪ algorythm,
78 # but since it is not the case we leave it as it is
79
80 # We plot again after pre-processing
81 d.cont <- melt(deposit[, c("age", "balance", "duration", "campaign",
  ↪ "previous")])
82 ggplot(d.cont, aes(x = value)) + facet_wrap(~variable, scales = "free") +
  ↪ geom_histogram() + theme(axis.text.x=element_text(angle=90, hjust=1,
  ↪ vjust=0.5))
83
84 # Now pdays is categorical
85 d.categ <- melt(deposit, measure.vars=c("job", "marital", "education",
  ↪ "housing", "loan", "contact", "default", "poutcome", "pdays"))
86 ggplot(d.categ, aes(x = value)) + facet_wrap(~variable, scales = "free") +
  ↪ geom_bar() + theme(axis.text.x=element_text(angle=90, hjust=1,
  ↪ vjust=0.5))
87
88 library(caret)
89 library(MASS)
90 library(e1071)
91 library(randomForest)
92
93 # PREPARING THE TRAINING AND TEST DATA
94 ## Since we want to use different methods, we need CV and a separate test
  ↪ set:
95
96 N <- nrow(deposit)
97 all.indexes <- 1:N
98
99 learn.indexes <- sample(1:N, round(2*N/3))
100 test.indexes <- all.indexes[-learn.indexes]
101
102 learn.data <- deposit[learn.indexes,]
103 original.learn.data <- original_data[learn.indexes,]

```

```

104 test.data <- deposit[test.indexes,]
105 original.test.data <- original_data[test.indexes,]
106
107 nlearn <- length(learn.indexes)
108 ntest <- N - nlearn
109
110 ##### MODELLING #####
111 ##### LOGISTIC REGRESSION #####
112 # We use Logistic Regression as recommended since it doesn't need a lot of
  ↳ preprocessing of the data and we also have a lot of categorical
  ↳ variables
113
114 # ORIGINAL DATA
115 # First aproximation with the original unchanged data & all variables
116 glm.fit = glm(subscribed~., data=original.learn.data, family="binomial")
117
118 # Observing the p-values, we can have an idea of the variables that have
  ↳ more importance in predicting our model,
119 # a low p-value indicates that we can reject the null hipotesis, thus that
  ↳ variable has an importance on our model,
120 # a higher p-value means that we can discard that variable
121 # so we can fit the mode again with just the variable that actually have
  ↳ an influence on our model
122 # We can discard the following since they affect our model less: age, job,
  ↳ marital, default, balance, pdays and previous
123 summary(glm.fit)
124
125 # We calculate the prediction with and without the discarded variables and
  ↳ compare the errors
126 glm.probs = predict(glm.fit, original.test.data, type="response")
127 glm.pred = rep("no", length(glm.probs))
128 glm.pred[glm.probs>.5] = "yes"
129
130 # We choose 3 values to represent our model performance: accuracy, error
  ↳ and precision, the last one is important
131 # because the bank wants to contact only those clients that are more
  ↳ probable to subscribe to the loan

```

```

132 # We can see that accuracy is high (90.07 %), but precission is low
    ↪ (34.15%), to solve this we lower the threshold
133 # for which a client may subscribe a loan (the probability) compare the
    ↪ values again, since for the bank clients
134 # with 30% probability of subscribing is probably worth to spend it's
    ↪ ressources contacting them
135 res.performance = table(glm.pred, original.test.data$subscribed)
136 res.accuracy = (res.performance[2,2] +
    ↪ res.performance[1,1])/sum(res.performance)*100
137 res.error = 100 - res.accuracy
138 res.precision = (res.performance[2,2])/(res.performance[2,2] +
    ↪ res.performance[1,2])*100
139
140 # Accuracy is slightly lower (90.03%) but precission has almost doubled
    ↪ (53.78%)
141 glm.pred[glm.probs>.3] = "yes"
142 res.performance = table(glm.pred, original.test.data$subscribed)
143 res.accuracy = (res.performance[2,2] +
    ↪ res.performance[1,1])/sum(res.performance)*100
144 res.error = 100 - res.accuracy
145 res.precision = (res.performance[2,2])/(res.performance[2,2] +
    ↪ res.performance[1,2])*100
146
147 # Now we fit the model without the variables that had less of an impact
148 glm.fit = glm(subscribed~.-age-job-marital-default-balance-pdays-previous,
    ↪ data=original.learn.data, family="binomial")
149 glm.probs = predict(glm.fit, original.test.data, type="response")
150 glm.pred = rep("no", length(glm.probs))
151
152 # The total accuracy decreases to 89.93% and precission to 53.32%, so
    ↪ using less variables makes our model a bit less accurate
153 # but the difference is really small so it's not really important to
    ↪ discard those variables
154 # If we have too many variables and computation time is important,
155 # we can also see that removing the ones we selected won't affect so much
    ↪ our model prediction
156 glm.pred[glm.probs>.3] = "yes"
157 res.performance = table(glm.pred, original.test.data$subscribed)

```

```

158 res.accuracy = (res.performance[2,2] +
  ↪ res.performance[1,1])/sum(res.performance)*100
159 res.error = 100 - res.accuracy
160 res.precision = (res.performance[2,2])/(res.performance[2,2] +
  ↪ res.performance[1,2])*100
161
162 # PREPROCESSED DATA
163 # We will fit with all the variables and also removing the ones that we
  ↪ mentioned before
164 glm.fit = glm(subscribed~., data=learn.data, family="binomial")
165 glm.probs = predict(glm.fit, test.data, type="response")
166 glm.pred = rep("no", length(glm.probs))
167
168 # Accuracy is 89.40% and precission is 57.95%, so our model is much more
  ↪ precise detecting clients
169 # that will probably buy the finantial product of the bank with the
  ↪ preprocessed data
170 glm.pred[glm.probs>.3] = "yes"
171 res.performance = table(glm.pred, original.test.data$subscribed)
172 res.accuracy = (res.performance[2,2] +
  ↪ res.performance[1,1])/sum(res.performance)*100
173 res.error = 100 - res.accuracy
174 res.precision = (res.performance[2,2])/(res.performance[2,2] +
  ↪ res.performance[1,2])*100
175
176 # Now we fit the model without the variables that had less of an impact
177 glm.fit = glm(subscribed~.-age-job-marital-default-balance-pdays-previous,
  ↪ data=learn.data, family="binomial")
178 glm.probs = predict(glm.fit, test.data, type="response")
179 glm.pred = rep("no", length(glm.probs))
180
181 # Accuracy: 89.47%, precision: 57.44%
182 # As before, there is a small reduction in accuracy and precision but the
  ↪ results with preprocessed data are better
183 glm.pred[glm.probs>.3] = "yes"
184 res.performance = table(glm.pred, original.test.data$subscribed)
185 res.accuracy = (res.performance[2,2] +
  ↪ res.performance[1,1])/sum(res.performance)*100

```

```

186 res.error = 100 - res.accuracy
187 res.precision = (res.performance[2,2])/(res.performance[2,2] +
    ↪ res.performance[1,2])*100
188
189 #### To get a better grasp at the performance of our model, we do k-fold
    ↪ cross validation
190 precision <- NULL
191 accuracy <- NULL
192 error <- NULL
193 k <- 100
194
195 # It may take a while to compute
196 for (i in 1:k)
197 {
198   N <- nrow(deposit)
199   all.indexes <- 1:N
200
201   # we choose 9/10s of the data as training data and the rest as test
    ↪ data
202   learn.indexes <- sample(1:N, round(9*N/10))
203   test.indexes <- all.indexes[-learn.indexes]
204
205   learn.data <- deposit[learn.indexes,]
206   test.data <- deposit[test.indexes,]
207
208   nlearn <- length(learn.indexes)
209   ntest <- N - nlearn
210   glm.fit = glm(subscribed ~ ., data=learn.data, family="binomial")
211   #glm.fit = glm(subscribed ~
    ↪ .-age-job-marital-default-balance-pdays-previous, data=learn.data,
    ↪ family="binomial")
212   glm.probs = predict(glm.fit, test.data, type="response")
213
214   glm.pred = rep("no", length(glm.probs))
215   glm.pred[glm.probs>.3] = "yes"
216
217   res.performance = table(glm.pred, test.data$subscribed)

```

```

218 accuracy[i] <- (res.performance[2,2] +
    ↪ res.performance[1,1])/sum(res.performance)*100
219 error[i] <- 100 - accuracy[i]
220 precision[i] <- (res.performance[2,2])/(res.performance[2,2] +
    ↪ res.performance[1,2])*100
221 }
222
223 # We can see that our model performs pretty well, even though the data is
    ↪ highly unbalanced
224 # Mean values with all the variables
225 # accuracy: 89.69%
226 # error: 10.31%
227 # precision: 58.80 %
228
229 # Mean values without the variables that influence less our model (swap
    ↪ the commented code in the previous bucle)
230 # accuracy: 89.74%
231 # error: 10.26%
232 # precision: 58.37 %
233 mean(accuracy)
234 mean(error)
235 mean(precision)
236
237 par(mfrow=c(1,3))
238 hist(accuracy, col='lightskyblue', border='lightskyblue4', xlab='Acuracy',
    ↪ main='Acuracy for CV', density=50)
239 hist(error, col='lightskyblue', border='lightskyblue4', xlab='Error',
    ↪ main='Error for CV', density=50)
240 hist(precision, col='lightskyblue', border='lightskyblue4',
    ↪ xlab='Precision', main='Precision for CV', density=50)
241
242 boxplot(accuracy, horizontal=T, col='lightskyblue',
    ↪ border='lightskyblue4', xlab='Acuracy', main='Acuracy for CV')
243 boxplot(error, horizontal=T, col='lightskyblue', border='lightskyblue4',
    ↪ xlab='Error', main='Error for CV')
244 boxplot(precision, horizontal=T, col='lightskyblue',
    ↪ border='lightskyblue4', xlab='Precision', main='Precision for CV')
245 dev.off()

```

```

246
247 # To compare the performance of our model we will also model with LDA and
    ↪ QDA and analyze their performances.
248 # Also we will test NaiveBayes and RandomForest
249 N <- nrow(deposit)
250 all.indexes <- 1:N
251
252 learn.indexes <- sample(1:N, round(2*N/3))
253 test.indexes <- all.indexes[-learn.indexes]
254
255 learn.data <- deposit[learn.indexes,]
256 test.data <- deposit[test.indexes,]
257
258 nlearn <- length(learn.indexes)
259 ntest <- N - nlearn
260
261 ##### LDA #####
262 # With LDA the precision is much lower, so we won't be using this model
263 # 10.73% error, 89.27% accuracy, 28.51% precision
264 lda.fit = lda(subscribed ~ ., data=learn.data)
265 lda.pred = predict(lda.fit, test.data)
266 lda.class = lda.pred$class
267
268 res.performance = table(lda.class, test.data$subscribed)
269 res.accuracy = (res.performance[2,2] +
    ↪ res.performance[1,1])/sum(res.performance)*100
270 res.error = 100 - res.accuracy
271 res.precision = (res.performance[2,2])/(res.performance[2,2] +
    ↪ res.performance[1,2])*100
272
273 ##### QDA #####
274 # Performs worse than LDA, but the precision is a bit higher so it detects
    ↪ better the subscriptions
275 # We confirm that both LDA and QDA are not suitable models to fit our
    ↪ problem
276 # 13.43% error, 86.57% accuracy, 37.99% precision
277 qda.fit <- qda(subscribed ~ ., data=learn.data)
278 qda.pred = predict(qda.fit, test.data)

```



```

279 qda.class = qda.pred$class
280
281 res.performance = table(qda.class, test.data$subscribed)
282 res.accuracy = (res.performance[2,2] +
  ↪ res.performance[1,1])/sum(res.performance)*100
283 res.error = 100 - res.accuracy
284 res.precision = (res.performance[2,2])/(res.performance[2,2] +
  ↪ res.performance[1,2])*100
285
286 ##### NAIVE BAYES #####
287 # It performs better than LDA and QDA, but worse than logistic regression
288 # 12.57% error, 87.43% accuracy, 43.89% precision
289 bayes.fit <- naiveBayes(subscribed ~ ., data = learn.data)
290 bayes.pred <- predict(bayes.fit, test.data)
291
292 res.performance = table(bayes.pred, test.data$subscribed)
293 res.accuracy = (res.performance[2,2] +
  ↪ res.performance[1,1])/sum(res.performance)*100
294 res.error = 100 - res.accuracy
295 res.precision = (res.performance[2,2])/(res.performance[2,2] +
  ↪ res.performance[1,2])*100
296
297 ##### Random Forest #####
298 # 9.30% error, 90.70% accuracy, 64.66% precision, so far the best method
299 rf <- randomForest(subscribed ~ ., data = original_data[learn.indexes,],
  ↪ ntree=300, proximity=FALSE)
300 rf.pred <- predict(rf, newdata=original_data[-learn.indexes,])
301
302 res.performance =
  ↪ table(Truth=original_data[-learn.indexes,]$subscribe,Pred=rf.pred)
303 res.accuracy = (res.performance[2,2] +
  ↪ res.performance[1,1])/sum(res.performance)*100
304 res.error = 100 - res.accuracy
305 res.precision = (res.performance[2,2])/(res.performance[2,2] +
  ↪ res.performance[1,2])*100
306
307 #### As with logistic regresion we do k-fold CV to confirm our model
  ↪ accuracy and precision

```

```

308 # We choose k=10 since randomForest has a high computation time
309 precision <- NULL
310 accuracy <- NULL
311 error <- NULL
312 k <- 10
313
314 # It may take quite a while to compute
315 for (i in 1:k)
316 {
317   N <- nrow(deposit)
318   all.indexes <- 1:N
319
320   # we choose 9/10s of the data as training data and the rest as test
321   ↪ data
322   learn.indexes <- sample(1:N, round(9*N/10))
323   test.indexes <- all.indexes[-learn.indexes]
324
325   nlearn <- length(learn.indexes)
326   ntest <- N - nlearn
327   #glm.fit = glm(subscribed ~ ., data=learn.data, family="binomial")
328   rf <- randomForest(subscribed ~ ., data = original_data[learn.indexes,],
329   ↪ ntree=300, proximity=FALSE)
330   rf.pred <- predict(rf, newdata=original_data[-learn.indexes,])
331
332   res.performance =
333   ↪ table(Truth=original_data[-learn.indexes,]$subscribe,Pred=rf.pred)
334   accuracy[i] <- (res.performance[2,2] +
335   ↪ res.performance[1,1])/sum(res.performance)*100
336   error[i] <- 100 - accuracy[i]
337   precision[i] <- (res.performance[2,2])/(res.performance[2,2] +
338   ↪ res.performance[1,2])*100
339 }
340
341 # Mean values
342 # accuracy: 90.83%
343 # error: 9.17 %
344 # precision: 64.66%
345 mean(accuracy)
346 mean(error)

```

```

341 mean(precision)
342
343 par(mfrow=c(1,3))
344 hist(accuracy, col='lightskyblue', border='lightskyblue4', xlab='Acuracy',
      ↪ main='Acuracy for CV', density=50)
345 hist(error, col='lightskyblue', border='lightskyblue4', xlab='Error',
      ↪ main='Error for CV', density=50)
346 hist(precision, col='lightskyblue', border='lightskyblue4',
      ↪ xlab='Precision', main='Precision for CV', density=50)
347
348 boxplot(accuracy, horizontal=T, col='lightskyblue',
      ↪ border='lightskyblue4', xlab='Acuracy', main='Acuracy for CV')
349 boxplot(error, horizontal=T, col='lightskyblue', border='lightskyblue4',
      ↪ xlab='Error', main='Error for CV')
350 boxplot(precision, horizontal=T, col='lightskyblue',
      ↪ border='lightskyblue4', xlab='Precision', main='Precision for CV')
351 dev.off()
352
353 # To conclude, random forest is the best method, followed by logistic
      ↪ regression, according to the results

```